

CSCI4180 Tutorial Week 9

Assignment 2

Parallel BFS Implementation

31 October 2013

Jeremy Chan
SHB118


cwchan@cse.cuhk.edu.hk

Overview of Assignment 2

- **Part 1 HBase (50%)**
 - Database-like system on Hadoop
 - Mostly configuration work
 - Will be covered next week in tutorial
- **Part 2 Parallel BFS (50%)**
 - Single-source shortest path lengths
 - Learn to run MapReduce **iteratively**
- **5% Bonus**


Definition

- Given a graph $G = (V, E)$ and a source node $v_s \in V$, find the **shortest path length** (distance) from v_s to every other **reachable** node in V .



First Step – Parsing Input

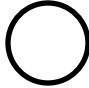
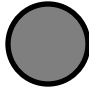

- From ***Input Format*** to ***adjacency list format***
 - So we make sure each node goes to the same mapper

| <u>Input Format</u> | | | <u>Adjacency List</u> | |
|---------------------|---|---|-----------------------|------|
| 1 | 3 |  | 1 | 3, 5 |
| 1 | 5 | | 3 | 1 |
| 3 | 1 | | 2 | 4 |
| 2 | 4 | | 4 | 2 |
| 4 | 2 | | | |

- You can design a separate Mapper/Reducer to do the transformation

Node Structure

- Three states (Color) of a node:

| | | |
|----------|-------------|---|
| – WHITE: | Not visited |  |
| – GRAY: | Discovered |  |
| – BLACK: | Processed |  |

- Stores distance from v_s as internal state

 Distance means the shortest path length

- Keeps track of neighbours by adjacency list

 Adjacency matrix is difficult for MapReduce

Hints for Designing the Node Class

- Use an enum for the color

```
public static enum Color {  
    WHITE, GRAY, BLACK  
};
```

- Use `int` for `id` and `distance`
- Use `List<Integer>` for the adjacency list
- Design a proper constructor
 - Think about the initial color, `id` and `distance`

How to Emit a Node?

The Java Way

- Implements a `toString()` method and a `fromString()` method in your Node Class
- Emit a node simply by a **`org.apache.hadoop.io.Text`**

IntWritable

e.g.

1

NodeID

Text

2,3,4,5 / 0 / BLACK

Adj. List

Distance

Color

- However, functions like `String.split()` could be slow

How to Emit a Node?

The Hadoop Way (Recommended)

- Implements ***Writable*** in the Node Class
 - You will need to provide implementation for the `readFields()` and `write()` method
- Emit a Node as the value directly
- Please refer to [Hadoop: The Definitive Guide](#), Chapter 4 - Hadoop I/O

Challenges in Parallel BFS

```
1: class MAPPER
```

```
2:   method MAP(nid  $n$ , node  $N$ )
```

```
3:      $d \leftarrow N.DISTANCE$ 
```

```
4:     EMIT(nid  $n$ ,  $N$ )
```

```
5:     for all nodeid  $m \in N.ADJACENCYLIST$  do
```

```
6:       EMIT(nid  $m$ ,  $d + 1$ )
```

```
1: class REDUCER
```

```
2:   method REDUCE(nid  $m$ , [ $d_1, d_2, \dots$ ])
```

```
3:      $d_{min} \leftarrow \infty$ 
```

```
4:      $M \leftarrow \emptyset$ 
```

```
5:     for all  $d \in \text{counts } [d_1, d_2, \dots]$  do
```

```
6:       if IsNode( $d$ ) then
```

```
7:          $M \leftarrow d$ 
```




```
8:       else if  $d < d_{min}$  then
```

```
9:          $d_{min} \leftarrow d$ 
```

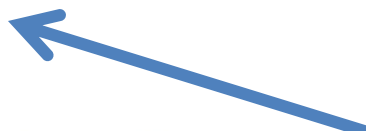
```
10:     $M.DISTANCE \leftarrow d_{min}$ 
```

```
11:    EMIT(nid  $m$ , node  $M$ )
```

How to process
GRAY nodes only?



How to emit two
kinds of data?




How to continue to
the next iteration?



Parallel BFS in MapReduce

- **MAPPER Design**



1. **Mappers** should not do anything to **BLACK** or **WHITE** nodes (not their turn to process), just emit the nodes “as-is”


2. **Mappers** processing **GRAY** node should

- 1) mark its neighbours as **GRAY**,
- 2) set neighbours' distance
- 3) mark itself **BLACK**
- 4) emits itself and its neighbours

Mapper's Limitation

1. **Mapper** which has v_1 emits v_3 as GRAY, another **Mapper** emits v_3 as WHITE

2. **Mapper** which has v_1 emits v_3 .distance = 2, another **Mapper** emits v_3 .distance = INF



3. **Mapper** which has v_1 emits v_3 without knowing its adjacency list, another **Mapper** which has v_3 emits v_3 with its adjacency list

Reducer solves consistency problem

Parallel BFS in MapReduce

- **REDUCER Design**
- For a Node, different mappers will provide conflicting information
- Think about how to combine the following three properties
 - Adjacency list
 - Distance
 - Color
- The emitted node can then be used for the Mapper in the next iteration

How to Chain MapReduce Jobs?

- Each “Map + Reduce” only explored one step further from v_s , we need to repeat this process
 - **Map1 → Reduce1 → Map2 → Reduce2 → Map3...**

The Java Way

- Drive jobs in the `main()` function using loops, counter and conditionals
- Different **Configuration** and **Job** Object in each iteration
- Set `job.waitForCompletion(true);`
- Jobs communicates by writing and reading intermediate files on Hadoop DFS
 - **For Job_i**
 - `FileInputFormat.addInputPath` (***OutputPath of Job_{i-1}***)

How to Chain MapReduce Jobs?

The Hadoop Way

- Use **org.apache.hadoop.mapred.jobcontrol**
- In this Hadoop version (0.20.203), use the following deprecated classes to chain jobs

`mapred.jobControl1.Jobcontrol1 ->`

`mapred.jobControl1.Job ->`

`mapred.jobConf ->`

`mapred.Mapper/Reducer/MapReduceBase`

How to Chain MapReduce Jobs?

- **Add a Job to JobControl**
 - `Job job1 = new Job (jobConf1)`
 - `Job job2 = new Job (jobConf2)`
- **Add dependency (e.g. 2 depends on 1)**
 - `job2.addDependingJob(job1)`
- **Add jobs to JobControl**
 - `JobControl jc = new JobControl ();`
 - `jc.addJob(job1);`
 - `jc.addJob(job2);`
- **Run JobControl**
 - `jc.run()`
- **Check if job is finished**
 - `jobControl.allFinished()` `// start this in a thread!`

When to Stop?

- Stop when all reachable nodes are processed
- **i.e. When there are no more GRAY nodes**
 - Nodes are either BLACK or WHITE
- How to check?
 - Use **org.apache.hadoop.mapred.Counters**

Hadoop Counter: Usage

1. Declare Counter

```
public static enum GrayCounter {  
    COUNT  
};
```

2. Increment Counter

```
context.getCounter(GrayCounter.COUNT).increment(1);
```

3. Retrieve Counter Value

```
int grayCount =  
job.getCounters().findCounter(ParallelBfs.GrayCounter.COUNT).getValue();
```

If you want to get the correct value of a Counter, only call `getValue()` when the job has finished running

Final Step – Clean-up

- Transform the output of the last run to the required format
- Only outputs the tuple which the node is *reachable*
- Feel free to use another set of Mapper/Reducer to do this step

Coming Next...

Nov 7

- HBase setup procedures
- How to write MapReduce programs for HBase

Due: Nov 14

- Demo on the next day
 - Only your last submission will be graded
 - Make sure your filenames are correct

Reminders

Tips

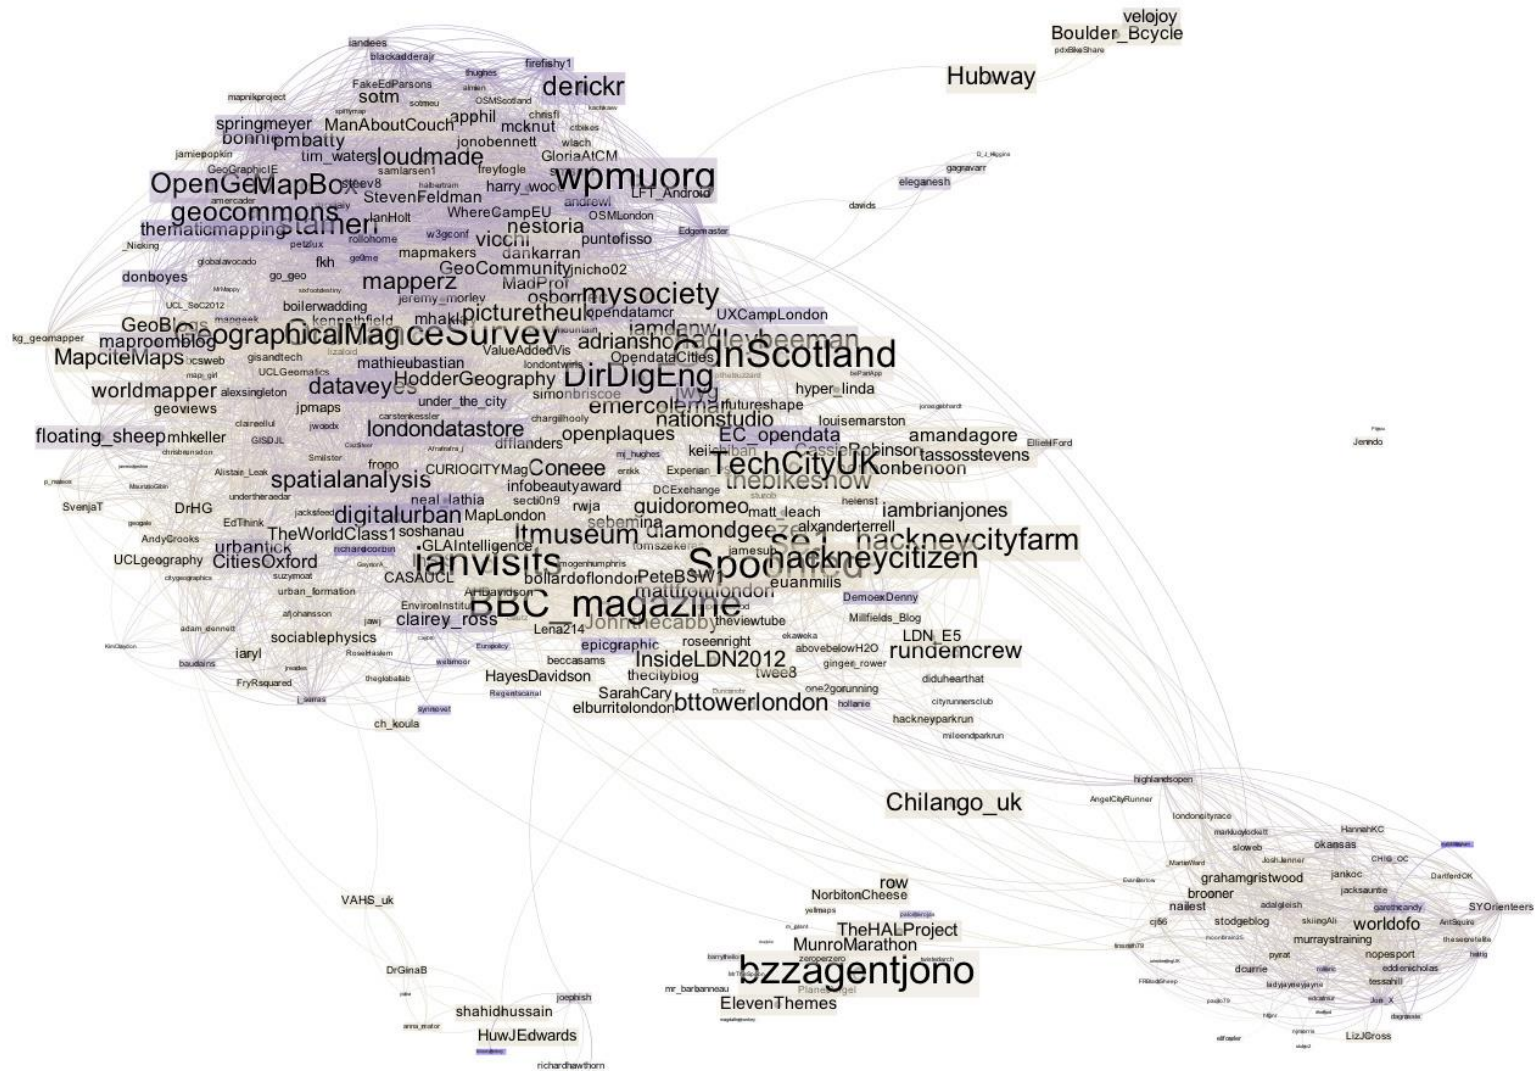
- Test your program correctness with hand-craft test cases (loops, directed edge, etc.)
- Running time depends on the number of passes of Map + Reduce
 - Largest test case in the demo will be Twitter_3, can take 5-15 mins to finish
- Should be more challenging than Assignment 1
- Please start early!

Questions?

Thank You

Visualizing Twitter Social Graph













Reference



<http://oliverobrien.co.uk/2012/07/six-degrees-of-twitter/>

BFS: Example






Reference



Credit: Dr. Matthew Tang's Data Structure Notes

BFS: Example

Reference



Q

- For the Assignment, the graph is **directed**
- Some nodes can be unreachable