

# **Lecture 9: Facebook Photo Storage**

CSCI4180 (Fall 2013)

Patrick P. C. Lee

# Outline

- Haystack: an efficient storage of billion of photos.
  - [http://static.usenix.org/event/osdi10/tech/full\\_papers/Beaver.pdf](http://static.usenix.org/event/osdi10/tech/full_papers/Beaver.pdf)
- Other Facebook attempts in building storage.
  - Facebook Messages.
    - Using HBase & Haystack.
  - Facebook Insight.
    - Using HBase for the storing results.

# Facebook Photo Storage

	April 2009	October 2010
Total	15 billion photos 60 billion images 1.5 petabytes ( $1.5 \times 10^{15}$ )	65 billion photos 260 billion images 20 petabytes ( $20 \times 10^{15}$ )
Upload Rate	220 million photos per week (25 terabytes per week)	1 billion photos per week (60 terabytes per week)
Serving Rate	550,000 images /sec	1 million images / sec

For each photo, four versions (called **images**) are stored and they are: **large**, **medium**, **small**, and **thumbnail**

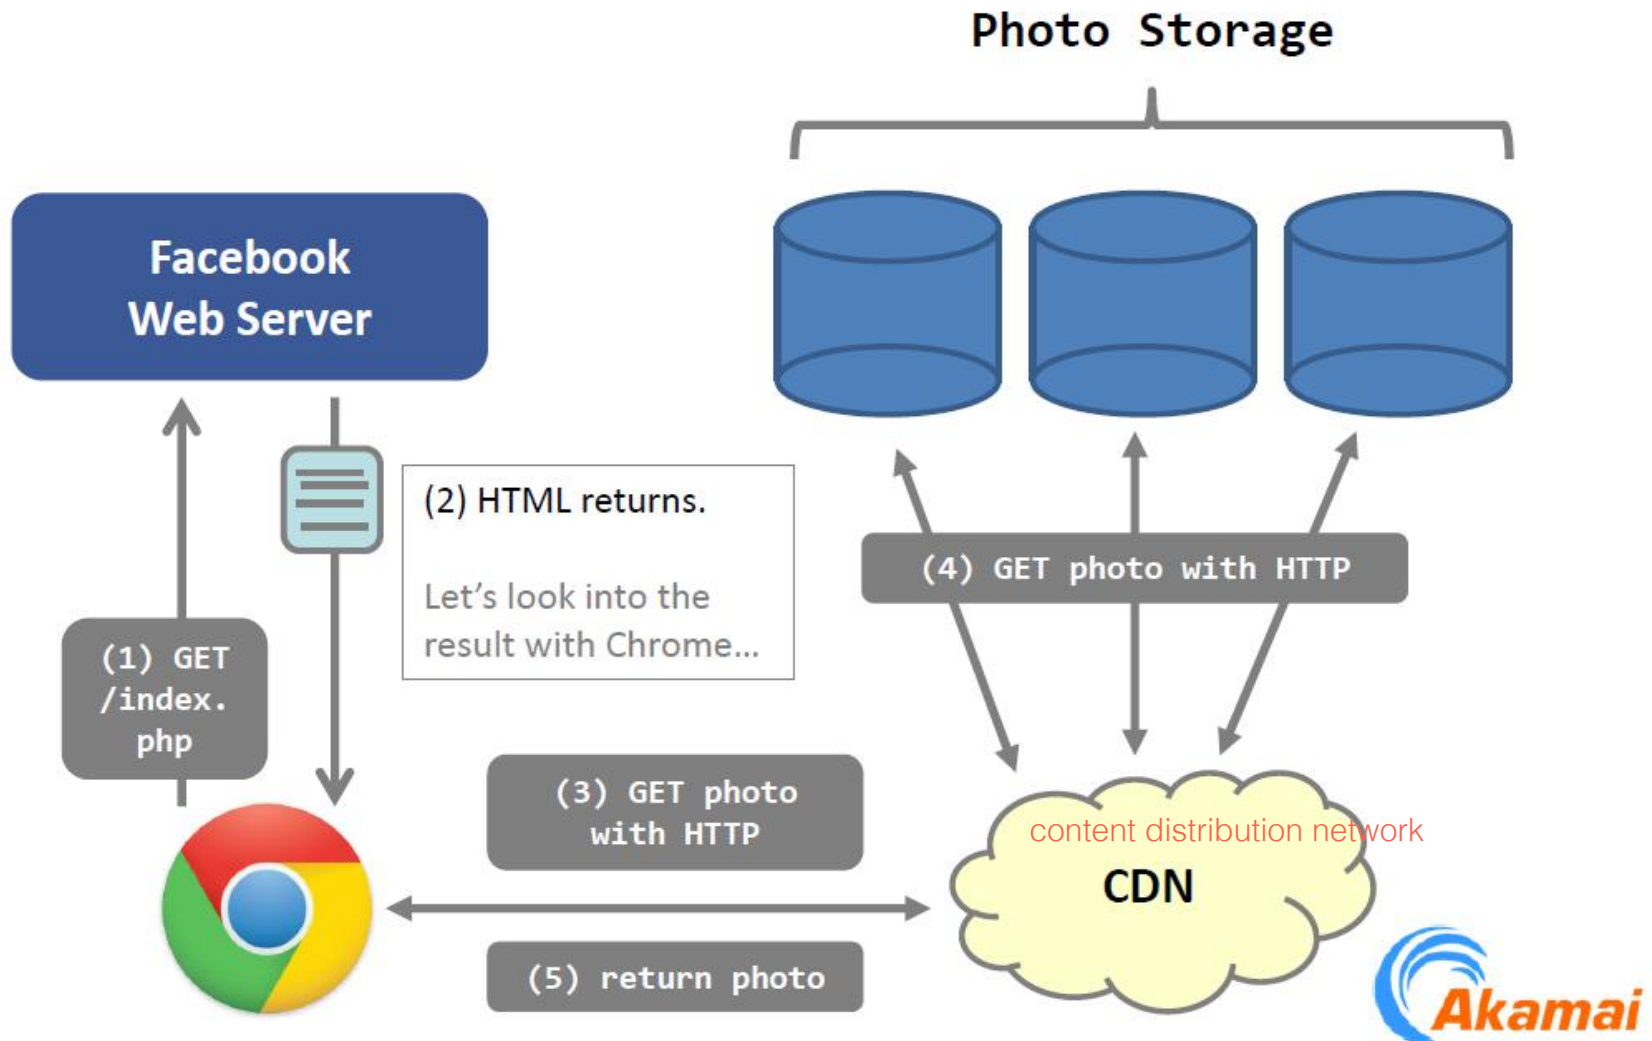
# Challenges

- Looking up a single file from a huge set of files...  
Similar to **finding a needle in haystack**
- Not a problem, if you understood FS design  
(from CSCI3150) *not scalable*
  - You need to know the pathname, i.e., the key.
  - The kernel digs out the metadata.
  - Size, block allocations, etc.
  - Return the file content
- Performance? That's the question!

# Design Goals

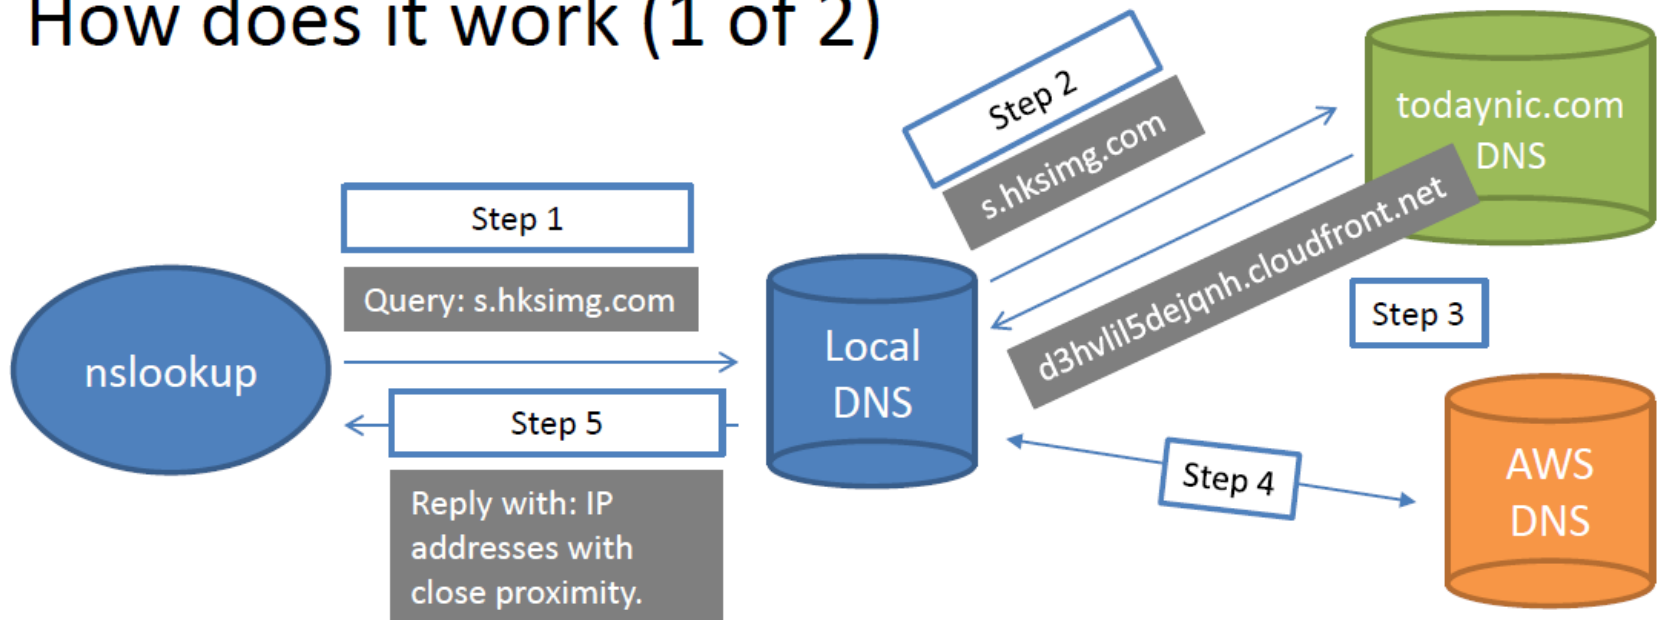
- High throughput and latency
  - Minimize disk I/Os
- Fault tolerant
  - Replicate photos in geographically distinct locations
- Cost effective
  - Cost per terabyte of usage must be kept low
- Simple
  - Straightforward to implement and maintain

# Typical Design



# Typical Design

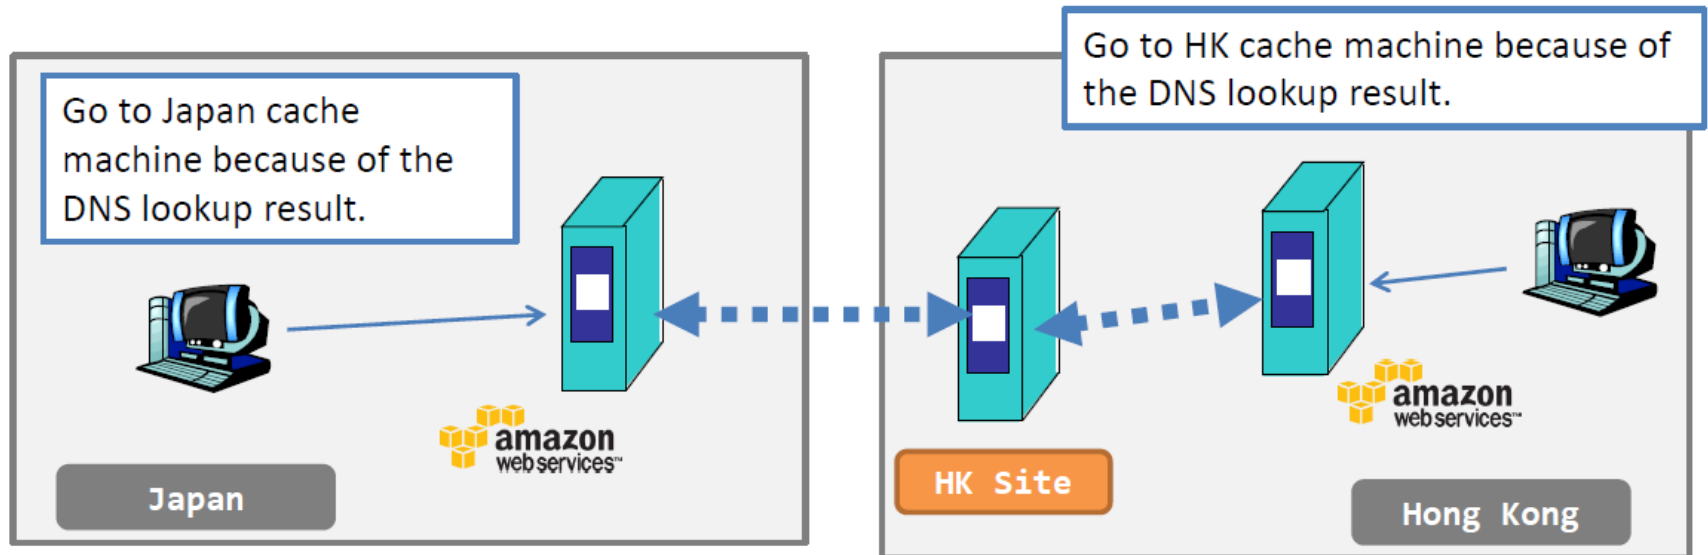
- CDN – **content delivery network** – is a network of nodes that **cache contents**.
- How does it work (1 of 2)



Example: CloudFront, <http://aws.amazon.com/cloudfront/>

# Typical Design

- CDN – **content delivery network** – is a network of nodes that **cache contents**.
- How does it work? (2 of 2)





# Typical Design

- CDN serves a large-scale cache
- Pros and Cons of CDN:
  - Pros: world-wide scale deployment, fast response time
  - Cons: control expiry of content
- Yet, CDN fits well typical web sites:
  - People access up-to-date data most of the time!
  - Old contents are seldom read.

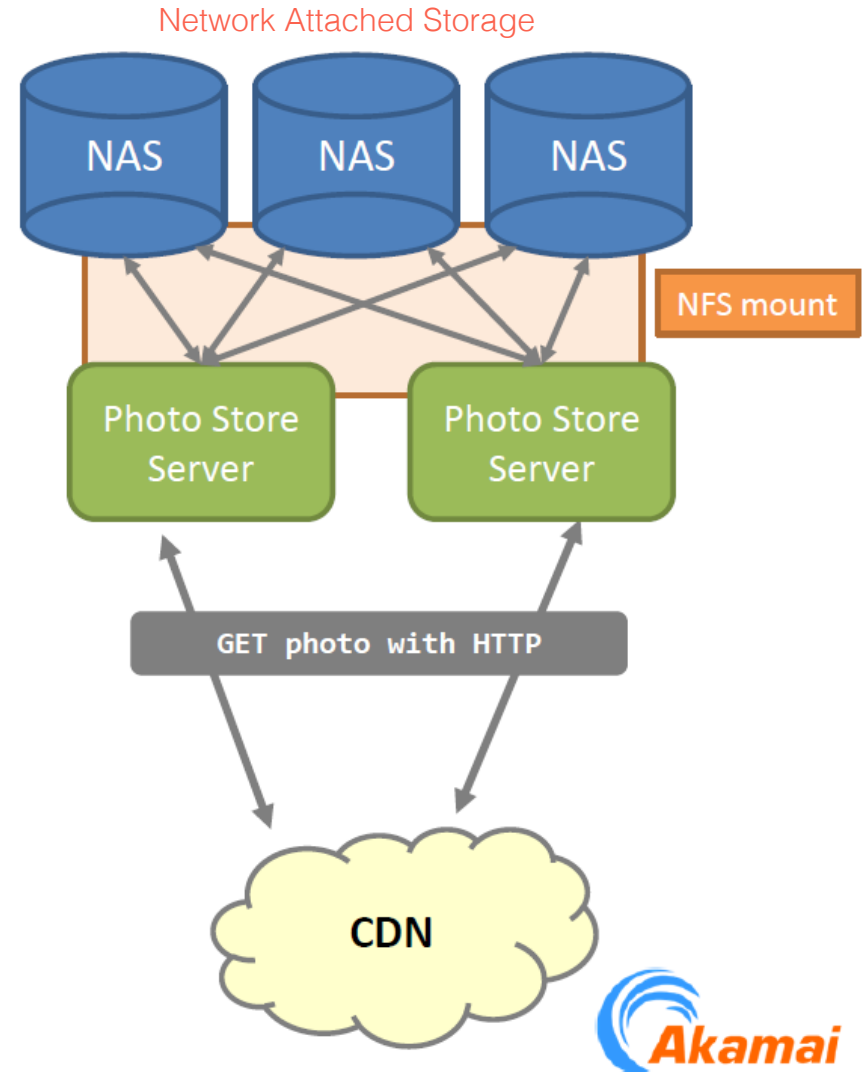
# Typical Design: Limitation

- Facebook (or social networking sites in general) has a very different access pattern from web sites
  - Old contents are frequently visited
  - Think about how you view your friend's album
- **Long tail** behavior
  - Requests from the long tail (or very old content) account of a significant amount of time

# (Old) NFS-based Design

➤ The old photo infrastructure had several tiers:

- **Upload tier** receives users' photo uploads, scales the original images and saves them on the NFS storage tier.
- **Photo serving tier** processes HTTP requests for images
- **NFS storage tier** built on top of commercial storage appliances.



# (Old) NFS-based Design

- An enormous amount of metadata is generated on the storage tier due to the namespace directories and file inodes.
- The amount of metadata far exceeds the caching abilities of the NFS storage tier
  - multiple disk I/Os per photo per upload/read request
- High degree of reliance on CDNs = expensive

# (Old) NFS-based Design

➤ With optimizations, we still need at least 3 I/O operations for each photo request:

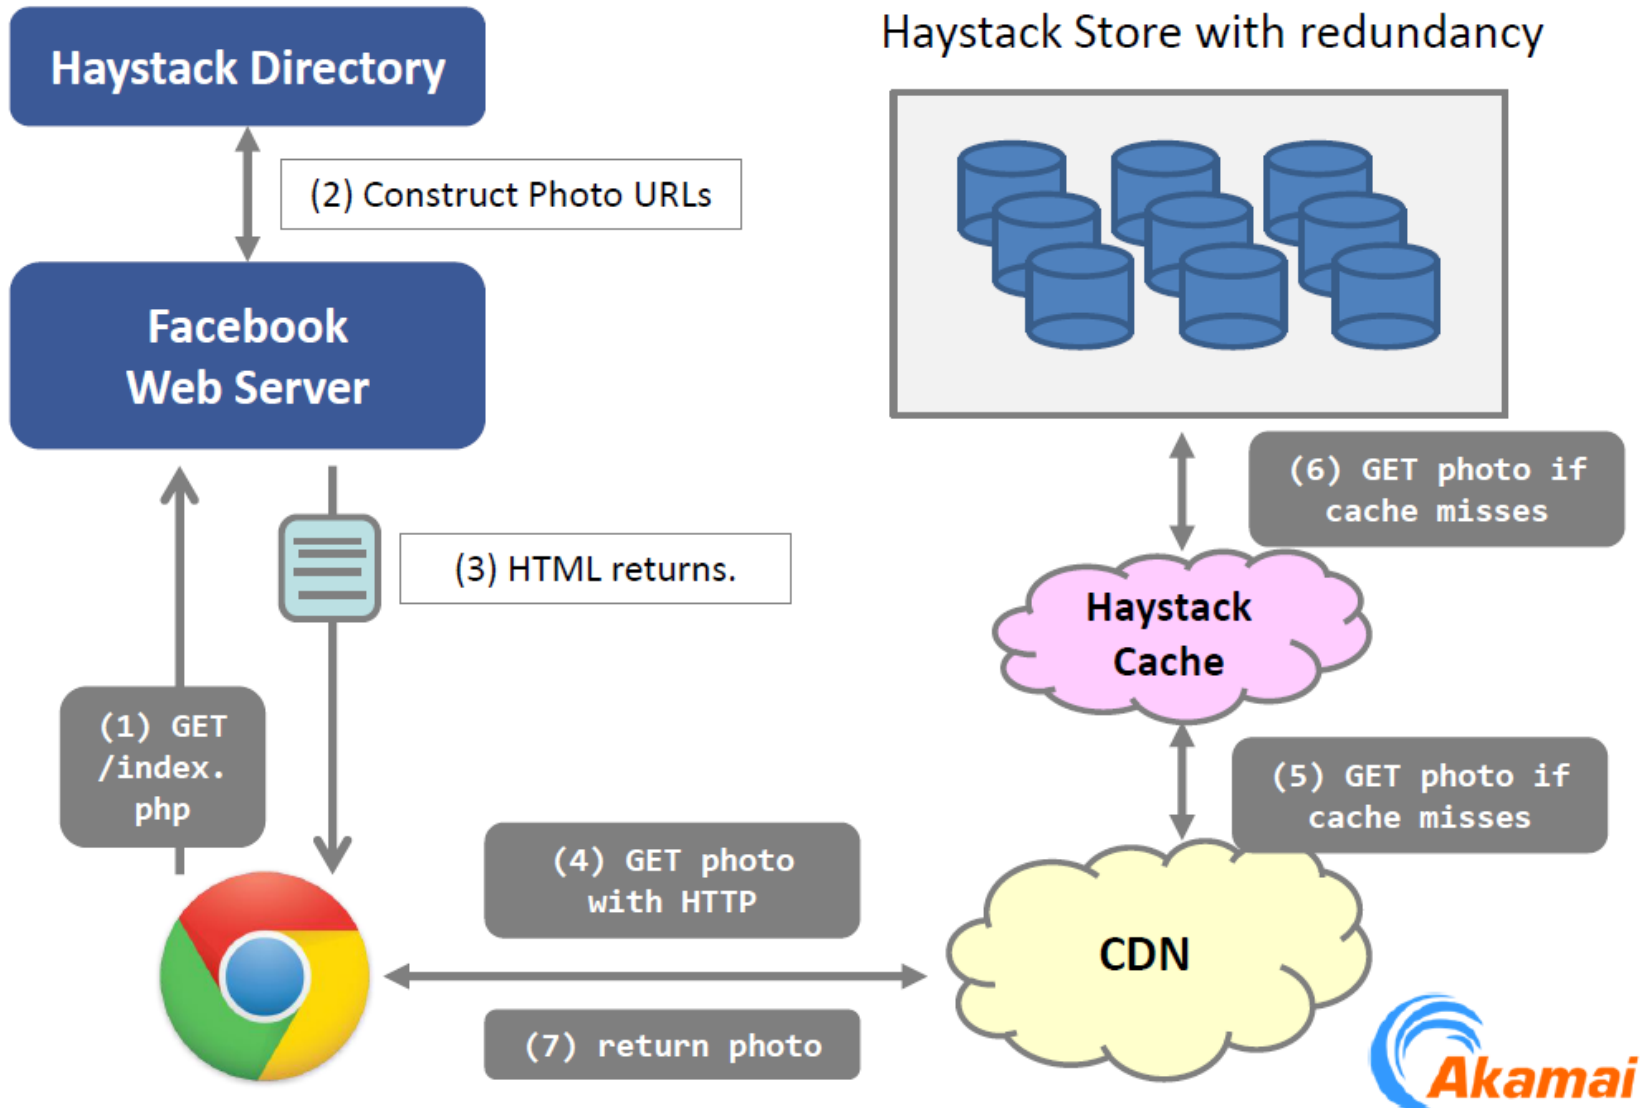
- Read the directory metadata into memory
- Load inode into memory
- Read file contents

too much I/O (4 versions of images for 1 photo), need to reduce the no. of I/O to save time

➤ Facebook's goal: further reduce I/Os

➤ Question: what is the minimum I/Os? 1 : read the file content

# Haystack Architecture



# Haystack: Key Components

- Develop customized file system with lightweight metadata operations
  - Haystack directory and Haystack store
- Offload CDN with a customized cache
  - Haystack cache

# User Operations

## ➤ User visits a page

- Web server forwards the request to the Directory
- The Directory constructs URL for each photo
  - `http://<CDN>/<Cache>/<Machine id>/<Logical volume, Photo>`
- CDN looks up the photo internally
  - If CDN lookup fails, the CDN strips the <CDN> part from the URL and contacts the Cache
  - If Cache lookup fails, the Cache strips the <Cache> part from the URL and contacts the Store



# Haystack Directory

## ➤ Four main functions:

- Provides a mapping from logical volumes to physical volumes
- Load balances writes across logical volumes
- Determines whether a photo request should be handled by the CDN or by the Haystack Cache
- Identifies logical volumes that are read-only
  - operational reasons
  - reached storage capacity

## ➤ Serves like the namenode in MapReduce

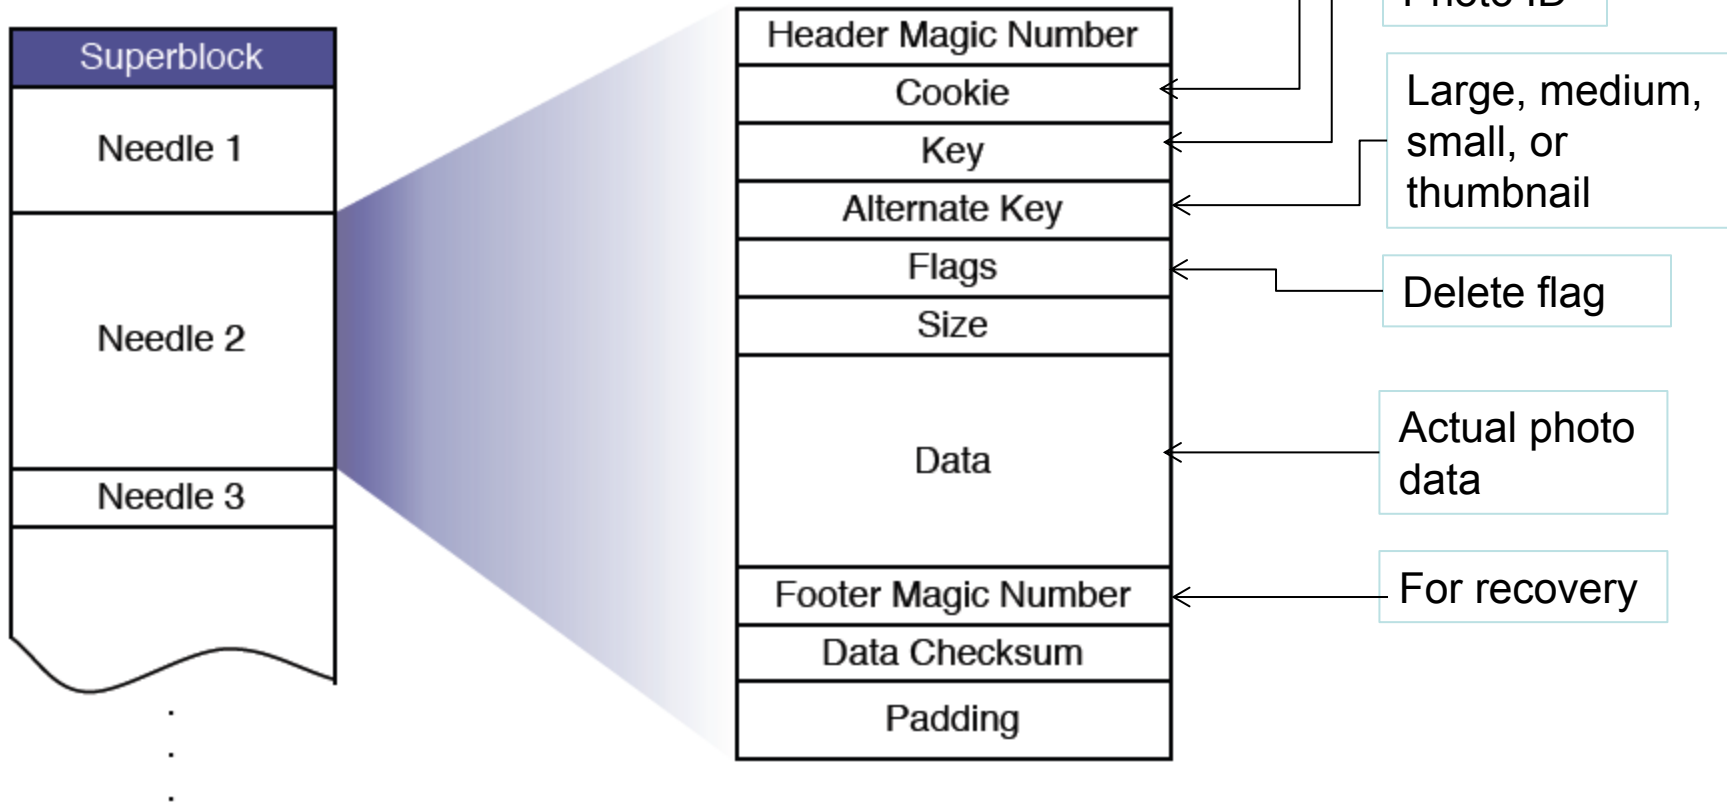
# Haystack Cache

- Distributed hash table, uses photo's id to locate cached data
- Receives HTTP requests from CDNs and browsers
  - If photo is in Cache, return the photo
  - If photo is not in Cache, fetches photo from the Store and returns the photo
- Add a photo to Cache if two conditions are met:
  - The request comes directly from a browser, not the CDN
  - The photo is fetched from a write-enabled Store machine

# Haystack Store

➤ A **Needle** represents an image

metadata, actual content



Layout of Haystack Store file


# Haystack Store

## ➤ Haystack store file

- A superblock file accompanied by many needles
- Log-structured format: append-only  
to reduce the seek

## ➤ Haystack index file

- provides minimal metadata required to locate a particular needle in the store
- Same order as store file
- A subset of store file, but without photo data
- For disaster recovery
- Less than 1% of store file



Layout of Haystack index file

# Haystack Store

- Each Store machine manages multiple physical volumes
- Can access a photo quickly using only:
  - the id of the corresponding logical volume
  - the file offset of the photo
- Handles three types of requests:
  - Read
  - Write
  - Delete

# Haystack Store: Read

- Each store machine maintains an in-memory index
- If both CDN and Haystack cache miss the photo:
  - Look up the in-memory index [ **0 I/O** ]
  - The offset in the Haystack store file is found.
  - Assume that the **store file is always opened**.
  - Seek and read [ **1 I/O** ]
    - Seek takes 1 I/O.
    - The succeeding read does not take any I/O ops since the seek operation has already positioned the disk head.
    - Since the photo is **sequentially written**, no further I/O ops incurred.
- •Totally, **1 I/O**

# Haystack Store: Write

- A multi-write operation: to **three** replicas
- Append 4 images to the haystack store file.
  - Seek to the end of file and sequentially write **[1 I/O]**.
- Then, append 4 images to the haystack index file.
  - Seek and then sequentially write **[ 1 I/O ]**.
- Update the in-memory index **[0 I/O]**.

# Haystack Store: Delete

- A multi-delete operation: to **three** replicas
- First, remove the files in the in-memory index.
  - This gives an illusion of fast file deletion.
  - Because any queries to the delete file would cause the in-memory index returning errors.
- Then, mark the 4 images of file deleted.
  - 4 seek-and-write.
- Garbage collection?
  - Create a copy of the store file
  - Skip deleted photos while copying
  - 25% of photos get deleted in a given year

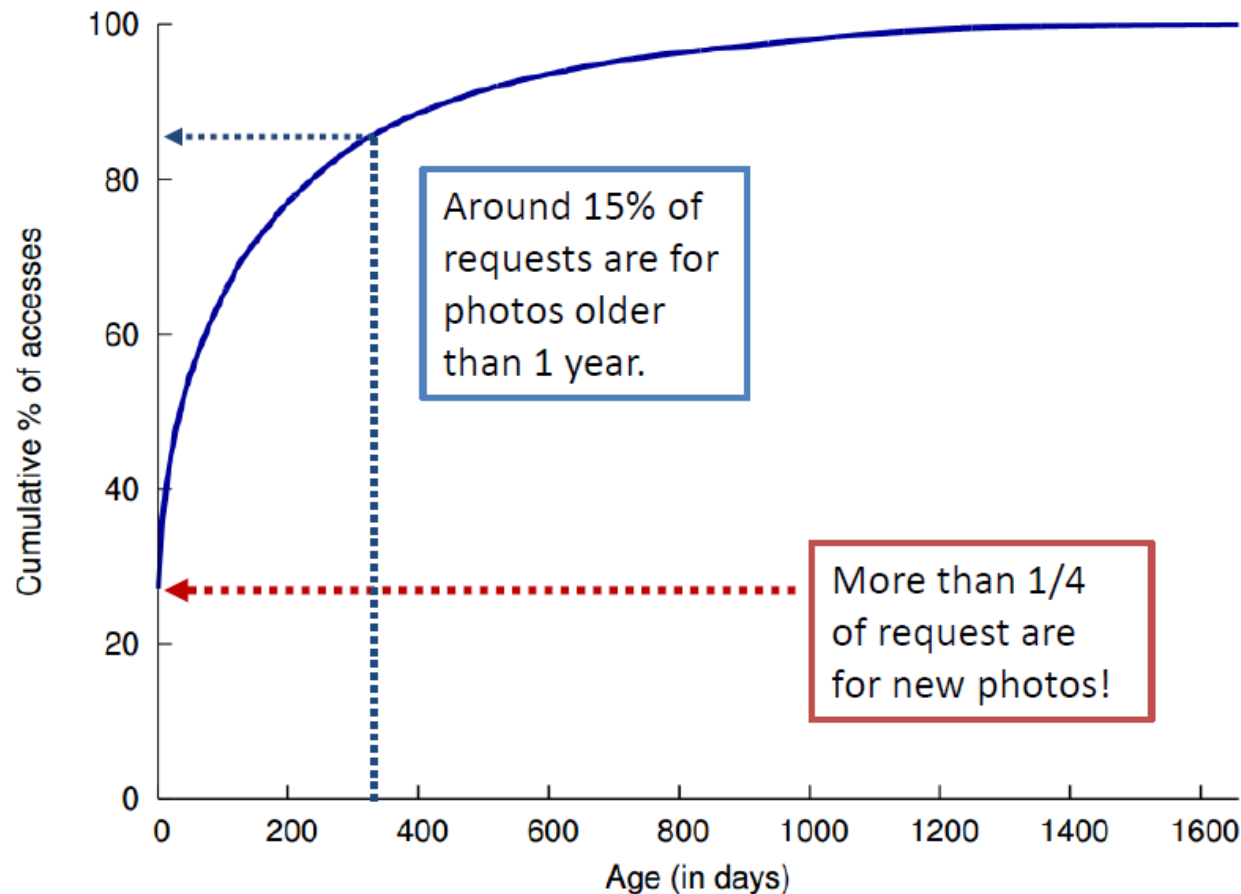


# Experiments

- Commodity machines running Haystack.
  - 2 x Quad-core CPUs.
  - 16GB – 32GB memory.
  - Hardware RAID controller.
  - 12+ 1TB SATA HDD
  
- Your machine can also be a Haystack

# Popularity

## ➤ Long tail behavior



# Cache Hit Rate




Figure 9: Cache hit rate for images that might be potentially stored in the Haystack Cache.

# Volume of Daily Total Traffic

4 different versions and each image is replicated to 3 servers.

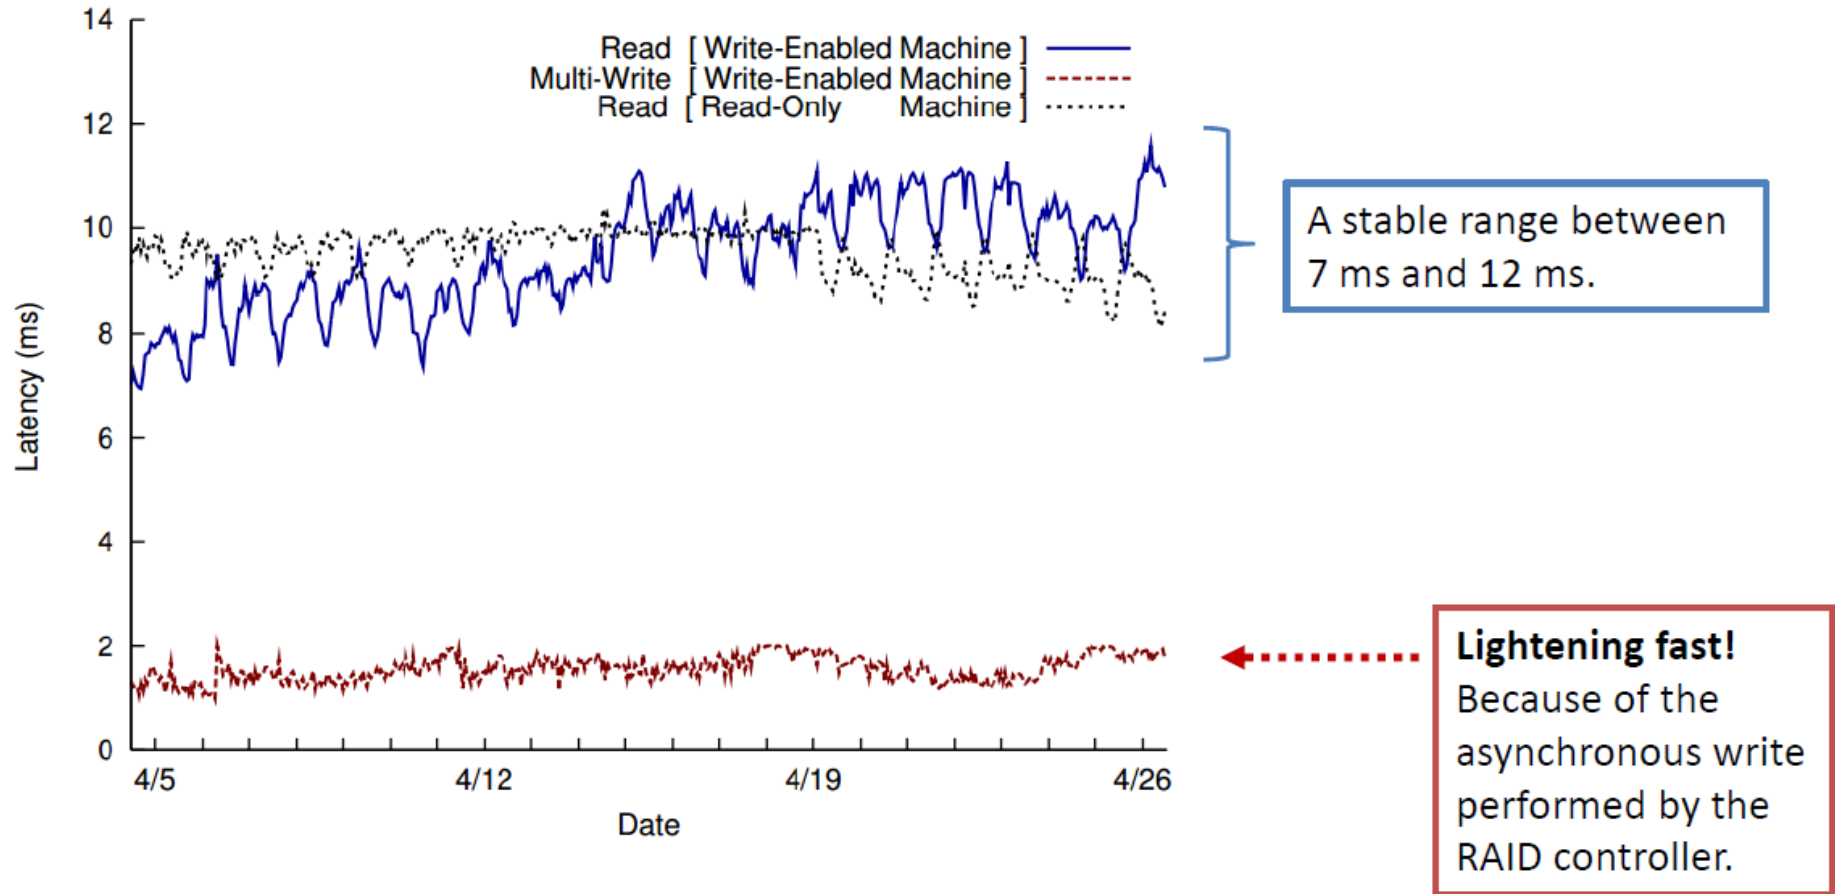
So, this matches Facebook's goal to build a custom system that saves io ops for small files.

Operations	Daily Counts
Photos Uploaded	~120 Million
Haystack Photos Written	~1.44 Billion
Photos Viewed	80-100 Billion
[ <i>Thumbnails</i> ]	10.2 %
[ <i>Small</i> ]	84.4 %
[ <i>Medium</i> ]	0.2 %
[ <i>Large</i> ]	5.2 %
Haystack Photos Read	10 Billion

Read is at least 700 times more than write.

Small photos in News feed!

# Duration of Workload: 3 weeks



# Haystack: Key Contributions

- Reduced disk I/O
  - 10 TB/node -> 10 GB of metadata
    - ~10 bytes per image, ~40 bytes per photo
    - This amount is easily cacheable!
- Simplified metadata
  - No directory structures/file names
    - 64-bit ID
  - Results in easier lookups
- Single photo serving and storage layer
  - Direct I/O path between client and storage
  - Results in higher bandwidth

# Haystack: Conclusions

- This is an efficient storage of billions of photos.
- What can we learn from Haystack?
  - Understand your problem!
    - Facebook problem is the haunting I/O requests over small photos.
  - Keep your solution simple.
    - Everybody who passed CSCI3150 will understand the merit of Haystack!

# Open Issues

- Is compaction (garbage collection) the best solution?  
Seems a bit expensive. Better ideas?
- What about album level abstraction?
  - Important/better if photos from the same album are placed sequentially or at least close together?
- Privacy concerns
  - Are cookies sufficient protection? Is there a better way?
  - What about security levels in Facebook? How are they enforced with respect to Haystack?
- How is consistency maintained between the Haystack and the CDN?
- Can we use erasure coding rather than replication?



# Inside Facebook

## ➤ Facebook messages

- Use **Apache HBase** stores
  - Small messages,
  - Message metadata, and
  - Search index of the users' messages.

## ➤ Facebook photos

- Use **Haystack** stores:
  - Large messages, and
  - Attachment.

➤ More: <http://www.facebook.com/Engineering>